

```
In [4]: ▶ import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from keras.datasets import mnist
from sklearn.metrics import accuracy_score

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Conv2D, Dense, MaxPooling2D
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.datasets import mnist
```

```
In [2]: ▶ #a. Loading and preprocessing the image data
(x_train,y_train),(x_test,y_test)=mnist.load_data()
```

```
In [3]: ▶ print(x_train.shape)
```

(60000, 28, 28)

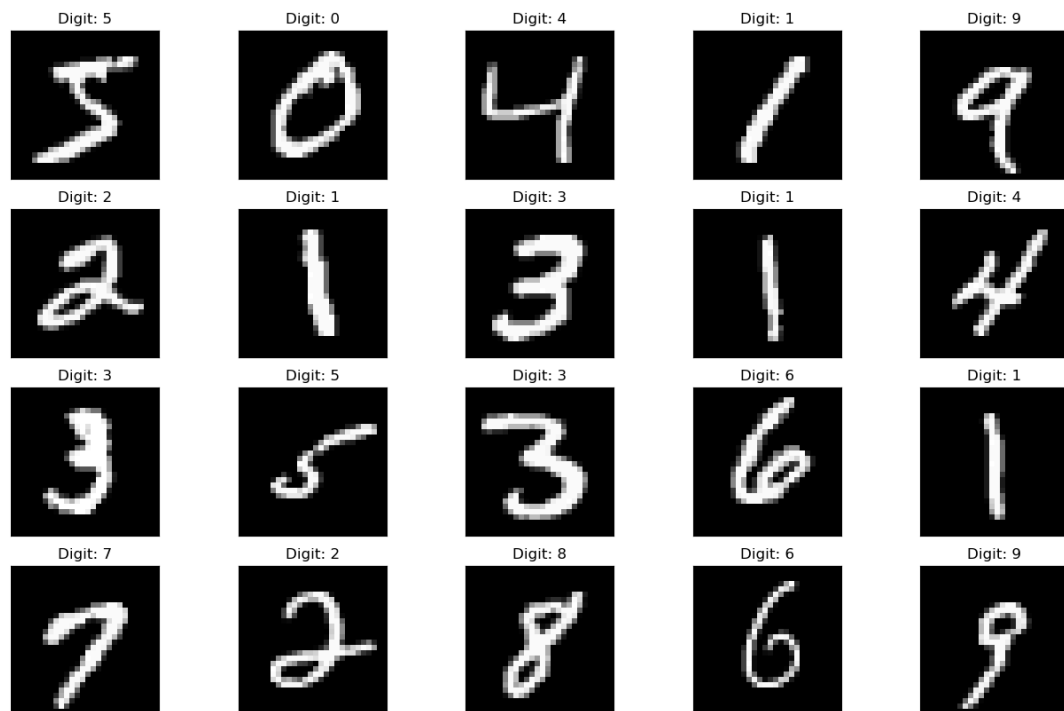
```
In [5]: ▶ x_train[0].min(),x_train[0].max()
```

Out[5]: (0, 255)

```
In [6]: ▶ x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
x_train[0].min(),x_train[0].max()
```

Out[6]: (0.0, 1.0)

```
In [7]: ▶ def plot_digit(image, digit, plt, i):
    plt.subplot(4, 5, i+1)
    plt.imshow(image, cmap=plt.get_cmap('gray'))
    plt.title(f"Digit: {digit}")
    plt.xticks([])
    plt.yticks([])
plt.figure(figsize=(16, 10))
for i in range(20):
    plot_digit(x_train[i], y_train[i], plt, i)
plt.show()
```



```
In [8]: ▶ x_train=x_train.reshape((x_train.shape + (1,)))
    x_test=x_test.reshape((x_test.shape + (1,)))
```

```
In [9]: ▶ y_train[0:20]
```

```
Out[9]: array([5, 0, 4, 1, 9, 2, 1, 3, 1, 4, 3, 5, 3, 6, 1, 7, 2, 8, 6, 9],
      dtype=uint8)
```

```
In [10]: #b. Defining the model's architecture
model = Sequential([
    Conv2D(32, (3, 3), activation="relu", input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(100, activation="relu"),
    Dense(10, activation="softmax")
])
```

WARNING:tensorflow:From C:\Users\shrey\AppData\Roaming\Python\Python311\site-packages\keras\src\backend.py:873: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

WARNING:tensorflow:From C:\Users\shrey\AppData\Roaming\Python\Python311\site-packages\keras\src\layers\pooling\max\_pooling2d.py:161: The name tf.nn.max\_pool is deprecated. Please use tf.nn.max\_pool2d instead.

```
In [11]: optimizer = SGD(learning_rate=0.01, momentum=0.9)
model.compile(
    optimizer=optimizer,
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
flatten (Flatten)	(None, 5408)	0
dense (Dense)	(None, 100)	540900
dense_1 (Dense)	(None, 10)	1010
=====		
Total params: 542230 (2.07 MB)		
Trainable params: 542230 (2.07 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [12]: ▶ #c. training the model
model.fit(x_train, y_train, epochs=10, batch_size=32)
```

Epoch 1/10

WARNING:tensorflow:From C:\Users\shrey\AppData\Roaming\Python\Python311\site-packages\keras\src\utils\tf\_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Users\shrey\AppData\Roaming\Python\Python311\site-packages\keras\src\engine\base\_layer\_utils.py:384: The name tf.executing\_eagerly\_outside\_functions is deprecated. Please use tf.compat.v1.executing\_eagerly\_outside\_functions instead.

1875/1875 [=====] - 21s 11ms/step - loss: 0.2422 - accuracy: 0.9283

Epoch 2/10

1875/1875 [=====] - 19s 10ms/step - loss: 0.0790 - accuracy: 0.9758

Epoch 3/10

1875/1875 [=====] - 20s 10ms/step - loss: 0.0483 - accuracy: 0.9854

Epoch 4/10

1875/1875 [=====] - 20s 11ms/step - loss: 0.0349 - accuracy: 0.9893

Epoch 5/10

1875/1875 [=====] - 20s 11ms/step - loss: 0.0255 - accuracy: 0.9921

Epoch 6/10

1875/1875 [=====] - 20s 11ms/step - loss: 0.0195 - accuracy: 0.9939

Epoch 7/10

1875/1875 [=====] - 20s 11ms/step - loss: 0.0140 - accuracy: 0.9959

Epoch 8/10

1875/1875 [=====] - 21s 11ms/step - loss: 0.0097 - accuracy: 0.9974

Epoch 9/10

1875/1875 [=====] - 20s 11ms/step - loss: 0.0074 - accuracy: 0.9980

Epoch 10/10

1875/1875 [=====] - 20s 11ms/step - loss: 0.0053 - accuracy: 0.9986


Out[12]: <keras.src.callbacks.History at 0x248369693d0>

```
In [14]: ▶ #d. Estimating the model's performance
predictions = np.argmax(model.predict(x_test), axis=-1)
accuracy_score(y_test, predictions)
```

313/313 [=====] - 1s 4ms/step

Out[14]: 0.9875

```
In [17]: ▶ import random
n=random.randint(0,9999)
plt.imshow(x_test[n])
plt.show()
predicted_value=model.predict(x_test)
print("Handwritten number in the image is=%d" %np.argmax(predicted_valu
```



313/313 [=====] - 1s 5ms/step  
Handwritten number in the image is=8

```
In [19]: ▶ x,y=model.evaluate(x_test,y_test, verbose=0)
print('Test loss:', x)
print('Test accuracy:', y)
```

Test loss: 0.04560748487710953  
Test accuracy: 0.987500011920929

```
In [ ]: ▶
```