


```
In [11]: ▶ #a. Import required libraries
import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, classification_report
import matplotlib.pyplot as plt
from tensorflow.keras import layers, models
```

```
In [28]: ▶ #b. Upload / access the dataset
dataset = pd.read_csv("creditcard.csv")
dataset
```

Out[28]:

	Time	V1	V2	V3	V4	V5	V6	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.2
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.0
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.7
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.2
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.5
...	...	...	...	...	...	...	...	...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.9
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.0
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.2
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.6
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.5

284807 rows × 31 columns



```
In [12]: ▶ # Preprocess the data (if needed)
# Normalize the data to have a mean of 0 and a standard deviation of 1
scaler = StandardScaler()
X = scaler.fit_transform(dataset.drop("Class", axis=1))
y = dataset["Class"]
```

```
In [13]: ▶ # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [14]: ▶ # Build and train the Autoencoder model
input_dim = X_train.shape[1]
```

```
In [15]: ▶ #c. Encoder converts it into latent representation
```

```
In [17]: # Encoder
encoder = models.Sequential([
    layers.Input(shape=(input_dim,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(16, activation='relu')
])
```

```
In [ ]: #d. Decoder networks convert it back to the original input
```

```
In [18]: # Decoder
decoder = models.Sequential([
    layers.Input(shape=(16,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(input_dim, activation='linear') # Using 'linear' activation
])

# Autoencoder
autoencoder = models.Sequential([
    encoder,
    decoder
])
```

```
In [ ]: #e . Compile the models with Optimizer, Loss, and Evaluation Metrics
```

```
In [19]: autoencoder.compile(optimizer='adam', loss='mean_squared_error')
autoencoder.fit(X_train, X_train, epochs=10, batch_size=32, shuffle=True)
```

WARNING:tensorflow:From C:\Users\shrey\AppData\Roaming\Python\Python311\site-packages\keras\src\optimizers\\_\_init\_\_.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

Epoch 1/10

WARNING:tensorflow:From C:\Users\shrey\AppData\Roaming\Python\Python311\site-packages\keras\src\utils\tf\_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

7121/7121 [=====] - 30s 4ms/step - loss: 0.3291 - val\_loss: 0.1998

Epoch 2/10

7121/7121 [=====] - 26s 4ms/step - loss: 0.1791 - val\_loss: 0.1571

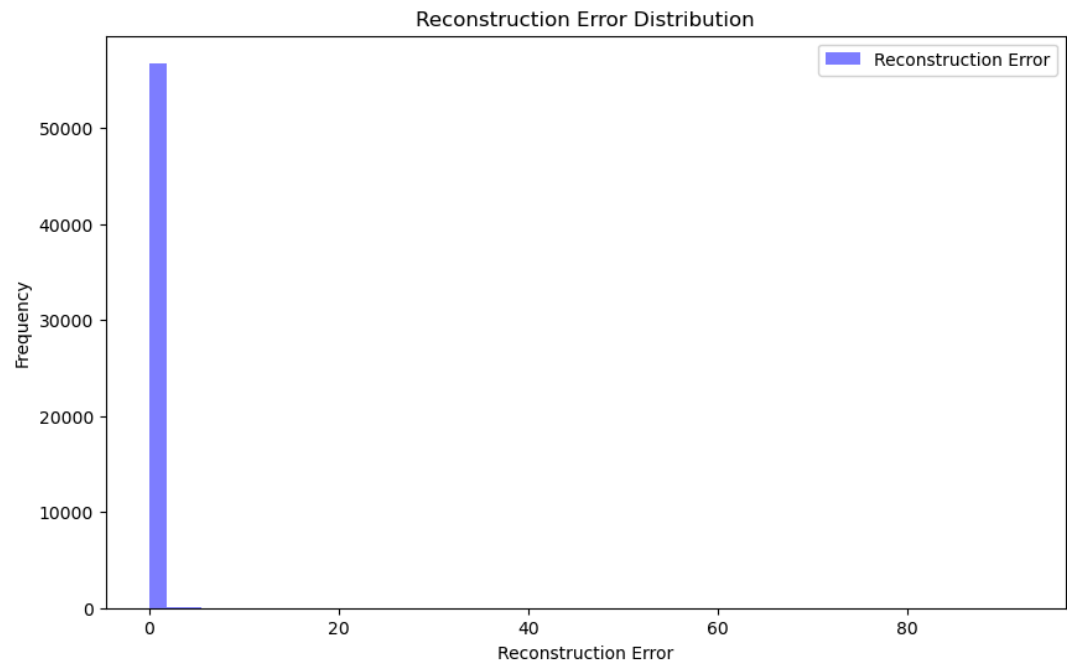
Epoch 3/10

7121/7121 [=====] - 26s 4ms/step - loss: 0.1481 - val\_loss: 0.1341

```
In [20]: # Detect anomalies and tune the threshold
y_pred = autoencoder.predict(X_test)
mse = np.mean(np.power(X_test - y_pred, 2), axis=1)
```

1781/1781 [=====] - 5s 2ms/step

```
In [21]: ▶ # Visualize the reconstruction error distribution
plt.figure(figsize=(10, 6))
plt.hist(mse, bins=50, alpha=0.5, color='b', label='Reconstruction Error')
plt.xlabel("Reconstruction Error")
plt.ylabel("Frequency")
plt.legend()
plt.title("Reconstruction Error Distribution")
plt.show()
```



```
In [22]: ▶ # Threshold tuning (iterate and adjust as needed)
thresholds = np.arange(0.1, 1.0, 0.1) # Adjust the step size as needed

for threshold in thresholds:
    anomalies = mse > threshold
```

```
In [30]: ▶ # Count the number of anomalies
num_anomalies = np.sum(anomalies)
print(f"Threshold: {threshold:.1f}, Number of anomalies: {num_anomalies}")
```

Threshold: 0.9, Number of anomalies: 558

```
In [24]: ▶ # Evaluate the model
print("Confusion Matrix:")
print(confusion_matrix(y_test, anomalies))

print("\nClassification Report:")
print(classification_report(y_test, anomalies))
```

Confusion Matrix:


```
[[56351  513]
 [   53   45]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	1.00	56864
1	0.08	0.46	0.14	98
accuracy			0.99	56962
macro avg	0.54	0.73	0.57	56962
weighted avg	1.00	0.99	0.99	56962

```
In [25]: ▶ import seaborn as sns
```

```
In [26]: ▶ plt.figure(figsize = (6, 4.75))
sns.heatmap(confusion_matrix(y_test, anomalies), annot = True, annot_kws={
plt.xticks([0.5, 1.5], rotation = 'horizontal')
plt.yticks([0.5, 1.5], rotation = 'horizontal')
plt.xlabel("Predicted label", fontsize = 14)
plt.ylabel("True label", fontsize = 14)
plt.title("Confusion Matrix", fontsize = 14)
plt.grid(False)
plt.show()
```



```
In [ ]: ▶
```